

20 def main(V, E, c, d):

1 children = find-children(V, E)
2 root = $V[0]$
3 result = find-flow($c, d, children, null, root$)
4 if result = null
5 return "no such flow"
6 else
7 return result

0 def find-flow($c, d, children, parent, root$):

1 flows = null
2 own-flow = $d[root]$
3 for each child $\in [children[root]]$:
4 child-flows = find-flow($c, d, children, root, child$)
5 if child-flows = null
6 return null
7 incoming = child-flows[(root, child)]
8 if [incoming] > $c[(root, child)]$
9 return null
10 own-flow = own-flow + incoming
11 flows = merge(flows, child-flows)
12 if parent ≠ null
13 flows[(parent, root)] = own-flow
14 return flows

Variables & Input & Functionsmain

V : a set of vertices^{in G}, labeled as letters (space permitting) labels can be arbitrary, letters are just easy.

E : a set of edges in G , each edge $e = (u, v)$ $u, v \in V$ is a tuple of 2 vertices

C : capacity dictionary $\{E \rightarrow \mathbb{R}^+\}$, for some $e \in E$, $c[e] = x$, where x is e 's cap.

d : demand dictionary $V \rightarrow \mathbb{R}$, for some $v \in V$, $d[v] = y$, where y is v 's demand

find-flow

find-children: this method takes a set of vertices V and a set of edges E from a tree G . It returns a dictionary $V \rightarrow \{V\}$ which maps a given vertex $v \in V$ to a list of its child vertices in tree G .

children: children dictionary $V \rightarrow \{V\}$, for some $v \in V$ ~~such that~~ $\text{children}[v] = \{c_1, c_2, c_3\}$ where $c_{1,2,3}$ are v 's children in G

parent: this is a vertex which is the parent of "root", null if root is root of G

root: root vertex for all recursive and external invocations of find-flow

Merge: this method merges 2 dictionaries, should not run into duplicate keys on merge. This isn't implemented, but can easily be implemented to run in $O(\min(s_1, s_2))$ where $\{s_1, s_2\}$ are the dictionary's sizes

Correctness

The algorithm traverses G in its entirety. By executing a depth first search, it recursively traverses ~~all~~ a path until it hits a leaf node, then returns flow values back up the call stack. Because of its recursive nature, no flow for a given node (or its edges) is calculated until such values are calculated for its children, ensuring that each flow value is correct. Line (8) of find-flow checks down-stream (incoming) flow against the appropriate edge's respective capacity to satisfy the capacity constraint.

Correctness (cont.'d)

aint. own-flow is initialized to root's demand. The algorithm then computes (recursively) all childrens' flows, updating own-flow to compute its own, assigning that flow to its parent if ~~the parent~~ is non-null. In the case that root is a leaf node, own-flow is initialized to root's demand, skips the for loop (as $\text{children}[\text{root}]$ is null for leaf nodes), and updates the flows object w/ its own-flow value before returning the flows object. This algorithm is also correct for the trivial cases $|V|=1, |E|=0$ and $|V|=2, |E|=1$.

Complexity

(find-flow)

The algorithm runs in $O(n)$ linear time where $n=|V|$. Because the algorithm executes a depth-first search (DFS), it visits each node precisely once, performing constant-time arithmetic and conditional logic per visit, keeping the complexity linear. *Find-children is not implemented here, as Prof. Spielman stipulated "If you like, you may assume that the tree is given to you in this [root-child-leaf] form." Anyways, the method can be implemented in $O(n \log n)$ with an initial sort and subsequent linear scan.